

# SCA-CNN: Spatial and Channel-wise Attention in Convolutional Networks for Image Captioning

Long Chen<sup>1</sup> Hanwang Zhang<sup>2</sup> Jun Xiao<sup>1\*</sup> Liqiang Nie<sup>3</sup> Jian Shao<sup>1</sup> Wei Liu<sup>4</sup> Tat-Seng Chua<sup>5</sup>

<sup>1</sup>Zhejiang University <sup>2</sup>Columbia University <sup>3</sup>Shandong University

<sup>4</sup>Tencent AI Lab <sup>5</sup>National University of Singapore

## Abstract

Visual attention has been successfully applied in structural prediction tasks such as visual captioning and question answering. Existing visual attention models are generally spatial, i.e., the attention is modeled as spatial probabilities that re-weight the last conv-layer feature map of a CNN encoding an input image. However, we argue that such spatial attention does not necessarily conform to the attention mechanism — a dynamic feature extractor that combines contextual fixations over time, as CNN features are naturally spatial, channel-wise and multi-layer. In this paper, we introduce a novel convolutional neural network dubbed SCA-CNN that incorporates Spatial and Channel-wise Attentions in a CNN. In the task of image captioning, SCA-CNN dynamically modulates the sentence generation context in multi-layer feature maps, encoding where (i.e., attentive spatial locations at multiple layers) and what (i.e., attentive channels) the visual attention is. We evaluate the proposed SCA-CNN architecture on three benchmark image captioning datasets: Flickr8K, Flickr30K, and MSCOCO. It is consistently observed that SCA-CNN significantly outperforms state-of-the-art visual attention-based image captioning methods.

## 1. Introduction

Visual attention has been shown effective in various structural prediction tasks such as image/video captioning [34, 36] and visual question answering [4, 35, 33]. Its success is mainly due to the reasonable assumption that human vision does not tend to process a whole image in its entirety at once; instead, one only focuses on selective parts of the whole visual space when and where as needed [5]. Specifically, rather than encoding an image into a static vector, attention allows the image feature to evolve from the

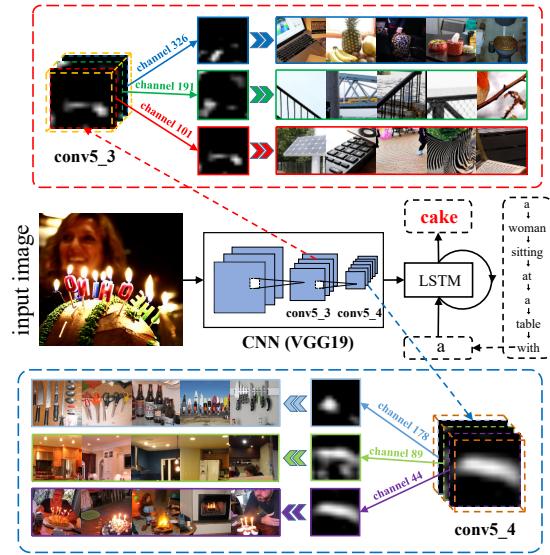


Figure 1. The illustration of channel-wise visual attention in two convolutional layers (*conv5\_3* and *conv5\_4* in VGG19) when predicting cake from the captioning a woman sitting at a table with cake. At each layer, top 3 attentive channels are visualized by showing the 5 most responsive receptive fields in the corresponding feature maps [40].

sentence context at hand, resulting in richer and longer descriptions for cluttered images. In this way, visual attention can be considered as a dynamic feature extraction mechanism that combines contextual fixations over time [19, 26].

State-of-the-art image features are generally extracted by deep Convolutional Neural Networks (CNNs) [8, 25, 32]. Starting from an input color image of the size  $W \times H \times 3$ , a convolutional layer consisting of  $C$ -channel filters scans the input image and output a  $W' \times H' \times C$  feature map, which will be the input for the next convolutional layer<sup>1</sup>. Each 2D slice of a 3D feature map encodes the spatial vi-

\*Corresponding author

<sup>1</sup>Each convolutional layer is optionally followed by a pooling, down-sampling, normalization, or a fully connected layer.

sual responses raised by a filter channel, where the filter performs as a pattern detector — lower-layer filters detect low-level visual cues like edges and corners while higher-level ones detect high-level semantic patterns like parts and object [40]. By stacking the layers, a CNN extracts image features through a hierarchy of visual abstractions. Therefore, CNN image features are essentially *spatial*, *channel-wise*, and *multi-layer*. However, most existing attention-based image captioning models only take into account the spatial characteristic [34], *i.e.*, those attention models merely modulate the sentence context into the last conv-layer feature map via spatially attentive weights.

In this paper, we will take full advantage of the three characteristics of CNN features for visual attention-based image captioning. In particular, we propose a novel Spatial and Channel-wise Attention-based Convolutional Neural Network, dubbed SCA-CNN, which learns to pay attention to every feature entry in the multi-layer 3D feature maps. Figure 1 illustrates the motivation of introducing channel-wise attention in multi-layer feature maps. First, since a channel-wise feature map is essentially a detector response map of the corresponding filter, channel-wise attention can be viewed as the process of selecting semantic attributes on the demand of the sentence context. For example, when we want to predict *cake*, our channel-wise attention (*e.g.*, in the *conv5\_3/conv5\_4* feature map) will assign more weights on channel-wise feature maps generated by filters according to the semantics like *cake*, *fire*, *light*, and *candle-like shapes*. Second, as a feature map is dependent on its lower-layer ones, it is natural to apply attention in multiple layers, so as to gain visual attention on multiple semantic abstractions. For example, it is beneficial to emphasize on lower-layer channels corresponding to more elemental shapes like *array* and *cylinder* that compose *cake*.

We validate the effectiveness of the proposed SCA-CNN on three well-known image captioning benchmarks: Flickr8K, Flickr30K and MSCOCO. SCA-CNN can significantly surpass the spatial attention model [34] by 4.8% in BLEU4. In summary, we propose a unified SCA-CNN framework to effectively integrate spatial, channel-wise, and multi-layer visual attention in CNN features for image captioning. In particular, a novel spatial and channel-wise attention model is proposed. This model is generic and thus can be applied to any layer in any CNN architecture such as popular VGG [25] and ResNet [8]. SCA-CNN helps us gain a better understanding of how CNN features evolve in the process of the sentence generation.

## 2. Related Work

We are interested in visual attention models used in the encoder-decoder framework for neural image/video captioning (NIC) and visual question answering (VQA), which

fall into the recent trend of connecting computer vision and natural language [14, 41, 24, 23, 42, 12]. Pioneering work on NIC [31, 13, 6, 30, 29] and VQA [1, 17, 7, 21] uses a CNN to encode an image or video into a static visual feature vector and then feed it into an RNN [9] to decode language sequences such as captions or answers.

However, the static vector does not allow the image feature adapting to the sentence context at hand. Inspired by the attention mechanism introduced in machine translation [2], where a decoder dynamically selects useful source language words or sub-sequence for the translation into a target language, visual attention models have been widely-used in NIC and VQA. We categorize these attention-based models into the following three domains that motivate our SCA-CNN:

- **Spatial Attention.** Xu *et al.* [34] proposed the first visual attention model in image captioning. In general, they used “hard” pooling that selects the most probably attentive region, or “soft” pooling that averages the spatial features with attentive weights. As for VQA, Zhu *et al.* [43] adopted the “soft” attention to merge image region features. To further refine the spatial attention, Yang *et al.* [35] and Xu *et al.* [33] applied a stacked spatial attention model, where the second attention is based on the attentive feature map modulated by the first one. Different from theirs, our multi-layer attention is applied on the multiple layers of a CNN. A common defect of the above spatial models is that they generally resort to weighted pooling on the attentive feature map. Thus, spatial information will be lost inevitably. More seriously, their attention is only applied in the last conv-layer, where the size of receptive field will be quite large and the differences between each receptive field region are quite limited, resulting in insignificant spatial attentions.

- **Semantic Attention.** Besides the spatial information, You *et al.* [37] proposed to select semantic concepts in NIC, where the image feature is a vector of confidences of attribute classifiers. Jia *et al.* [11] exploited the correlation between images and their captions as the global semantic information to guide the LSTM generating sentences. However, these models require external resources to train these semantic attributes. In SCA-CNN, each filter kernel of a convolutional layer servers as a semantic detectors [40]. Therefore, the channel-wise attention of SCA-CNN is similar to semantic attention.

- **Multi-layer Attention.** According to the nature of CNN architecture, the sizes of respective fields corresponding to different feature map layers are different. To overcome the weakness of large respective field size in the last conv-layer attention, Seo *et al.* [22] proposed a multi-layer attention networks. In compared with theirs, SCA-

CNN also incorporates the channel-wise attention at multiple layers.

### 3. Spatial and Channel-wise Attention CNN

#### 3.1. Overview

We adopt the popular encoder-decoder framework for image caption generation, where a CNN first encodes an input image into a vector and then an LSTM decodes the vector into a sequence of words. As illustrated in Figure 2, SCA-CNN makes the original CNN multi-layer feature maps adaptive to the sentence context through channel-wise attention and spatial attention at multiple layers.

Formally, suppose that we want to generate the  $t$ -th word of the image caption. At hand, we have the last sentence context encoded in the LSTM memory  $\mathbf{h}_{t-1} \in \mathbb{R}^d$ , where  $d$  is the hidden state dimension. At the  $l$ -th layer, the spatial and channel-wise attention weights  $\gamma^l$  are a function of  $\mathbf{h}_{t-1}$  and the current CNN features  $\mathbf{V}^l$ . Thus, SCA-CNN modulates  $\mathbf{V}^l$  using the attention weights  $\gamma^l$  in a recurrent and multi-layer fashion as:

$$\begin{aligned}\mathbf{V}^l &= \text{CNN}(\mathbf{X}^{l-1}), \\ \gamma^l &= \Phi(\mathbf{h}_{t-1}, \mathbf{V}^l), \\ \mathbf{X}^l &= f(\mathbf{V}^l, \gamma^l).\end{aligned}\quad (1)$$

where  $\mathbf{X}^l$  is the modulated feature,  $\Phi(\cdot)$  is the spatial and channel-wise attention function that will be detailed in Section 3.2 and 3.3,  $\mathbf{V}^l$  is the feature map output from previous conv-layer, *e.g.*, convolution followed by pooling, down-sampling or convolution [25, 8], and  $f(\cdot)$  is a linear weighting function that modulates CNN features and attention weights. Different from existing popular modulating strategy that sums up all visual features based on attention weights [34], function  $f(\cdot)$  applies element-wise multiplication. So far, we are ready to generate the  $t$ -th word by:

$$\begin{aligned}\mathbf{h}_t &= \text{LSTM}(\mathbf{h}_{t-1}, \mathbf{X}^L, y_{t-1}), \\ y_t &\sim p_t = \text{softmax}(\mathbf{h}_t, y_{t-1}).\end{aligned}\quad (2)$$

where  $L$  is the total number of conv-layers;  $p_t \in \mathbb{R}^{|\mathcal{D}|}$  is a probability vector and  $\mathcal{D}$  is a predefined dictionary including all caption words.

Note that  $\gamma^l$  is of the same size as  $\mathbf{V}^l$  or  $\mathbf{X}^l$ , *i.e.*,  $W^l \times H^l \times C^l$ . It will require  $\mathcal{O}(W^l H^l C^l k)$  space for attention computation, where  $k$  is the common mapping space dimension of CNN feature  $\mathbf{V}^l$  and hidden state  $\mathbf{h}_{t-1}$ . It is prohibitively expensive for GPU memory when the feature map size is so large. Therefore, we propose an approximation that learns spatial attention weights  $\alpha^l$  and channel-wise attention weights  $\beta^l$  separately:

$$\alpha^l = \Phi_s(\mathbf{h}_{t-1}, \mathbf{V}^l), \quad (3)$$

$$\beta^l = \Phi_c(\mathbf{h}_{t-1}, \mathbf{V}^l). \quad (4)$$

Where  $\Phi_c$  and  $\Phi_s$  represent channel-wise and spatial attention model respectively. This will greatly reduce the memory cost into  $\mathcal{O}(W^l H^l k)$  for spatial attention and  $\mathcal{O}(C^l k)$  for channel-wise attention, respectively.

#### 3.2. Spatial Attention

In general, a caption word only relates to partial regions of an image. For example, in Figure 1, when we want to predict *cake*, only image regions which contain cake are useful. Therefore, applying a global image feature vector to generate caption may lead to sub-optimal results due to the irrelevant regions. Instead of considering each image region equally, spatial attention mechanism attempts to pay more attention to the semantic-related regions. Without loss of generality, we discard the layer-wise superscript  $l$ . We reshape  $\mathbf{V} = [\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_m]$  by flattening the width and height of the original  $\mathbf{V}$ , where  $\mathbf{v}_i \in \mathbb{R}^C$  and  $m = W \cdot H$ . We can consider  $\mathbf{v}_i$  as the visual feature of the  $i$ -th location. Given the previous time step LSTM hidden state  $\mathbf{h}_{t-1}$ , we use a single-layer neural network followed by a softmax function to generate the attention distributions  $\alpha$  over the image regions. Below are the definitions of the spatial attention model  $\Phi_s$ :

$$\begin{aligned}\mathbf{a} &= \tanh((\mathbf{W}_s \mathbf{V} + b_s) \oplus \mathbf{W}_{hs} \mathbf{h}_{t-1}), \\ \alpha &= \text{softmax}(\mathbf{W}_i \mathbf{a} + b_i).\end{aligned}\quad (5)$$

where  $\mathbf{W}_s \in \mathbb{R}^{k \times C}$ ,  $\mathbf{W}_{hs} \in \mathbb{R}^{k \times d}$ ,  $\mathbf{W}_i \in \mathbb{R}^k$  are transformation matrices that map image visual features and hidden state to a same dimension. We denote  $\oplus$  as the addition of a matrix and a vector. And the addition between a matrix and a vector is performed by adding each column of the matrix by the vector.  $b_s \in \mathbb{R}^k$ ,  $b_i \in \mathbb{R}^1$  are model biases.

#### 3.3. Channel-wise Attention

Note that the spatial attention function in Eq (3) still requires the visual feature  $\mathbf{V}$  to calculate the spatial attention weights, but the visual feature  $\mathbf{V}$  used in spatial attention is in fact not attention-based. Hence, we introduce a channel-wise attention mechanism to attend the features  $\mathbf{V}$ . It is worth noting that each CNN filter performs as a pattern detector, and each channel of a feature map in CNN is a response activation of the corresponding convolutional filter. Therefore, applying an attention mechanism in channel-wise manner can be viewed as a process of selecting semantic attributes.

For channel-wise attention, we first reshape  $\mathbf{V}$  to  $\mathbf{U}$ , and  $\mathbf{U} = [\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_C]$ , where  $\mathbf{u}_i \in \mathbb{R}^{W \times H}$  represents the  $i$ -th channel of the feature map  $\mathbf{V}$ , and  $C$  is the total number of channels. Then, we apply mean pooling for each channel to obtain the channel feature  $\mathbf{v}$ :

$$\mathbf{v} = [v_1, v_2, \dots, v_C], \mathbf{v} \in \mathbb{R}^C, \quad (6)$$

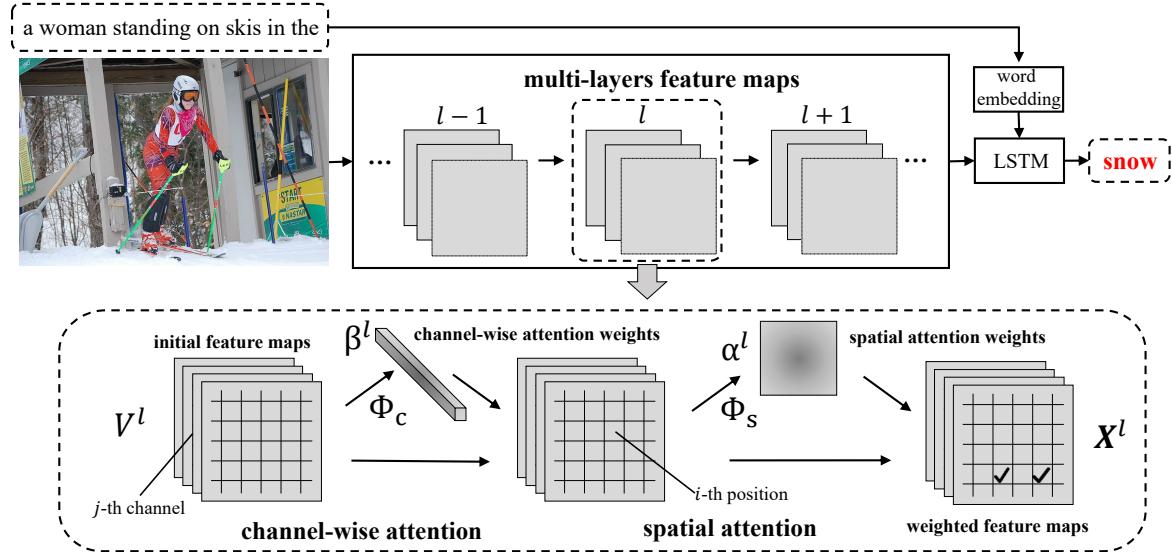


Figure 2. The overview of our proposed SCA-CNN. For the  $l$ -th layer, initial feature map  $\mathbf{V}^l$  is the output of  $(l - 1)$ -th conv-layer. We first use the channel-wise attention function  $\Phi_c$  to obtain the channel-wise attention weights  $\beta^l$ , which are multiplied in channel-wise of the feature map. Then, we use the spatial attention function  $\Phi_s$  to obtain the spatial attention weights  $\alpha^l$ , which are multiplied in each spatial regions, resulting in an attentive feature map  $\mathbf{X}^l$ . Different orders of two attention mechanism are discussed in Section 3.3.

where scalar  $v_i$  is the mean of vector  $\mathbf{u}_i$ , which represents the  $i$ -th channel features. Following the definition of the spatial attention model, the channel-wise attention model  $\Phi_c$  can be defined as follows:

$$\begin{aligned}\mathbf{b} &= \tanh((\mathbf{W}_c \otimes \mathbf{v} + b_c) \oplus \mathbf{W}_{hc} \mathbf{h}_{t-1}), \\ \beta &= \text{softmax}(\mathbf{W}'_i \mathbf{b} + b'_i).\end{aligned}\quad (7)$$

where  $\mathbf{W}_c \in \mathbb{R}^k$ ,  $\mathbf{W}_{hc} \in \mathbb{R}^{k \times d}$ ,  $\mathbf{W}'_i \in \mathbb{R}^k$  are transformation matrices,  $\otimes$  represents the outer product of vectors.  $b_c \in \mathbb{R}^k$ ,  $b'_i \in \mathbb{R}^1$  are bias terms.

According to different implementation order of channel-wise attention and spatial attention, there exists two types of model which incorporating both two attention mechanisms. We distinguish between the two types as follows:

**Channel-Spatial.** The first type dubbed Channel-Spatial (C-S) applies channel-wise attention before spatial attention. The flow chart of C-S type is illustrated in Figure 2. At first, given an initial feature map  $\mathbf{V}$ , we adopt channel-wise attention  $\Phi_c$  to obtain the channel-wise attention weights  $\beta$ . Through a linear combination of  $\beta$  and  $\mathbf{V}$ , we obtain a channel-wise weighted feature map. Then we feed the channel-wise weighted feature map to the spatial attention model  $\Phi_s$  and obtain the spatial attention weights  $\alpha$ . After attaining two attention weights  $\alpha$  and  $\beta$ , we can feed  $\mathbf{V}, \beta, \alpha$  to modulate function  $f$  to calculate the modulated

feature map  $\mathbf{X}$ . All processes are summarized as follows:

$$\begin{aligned}\beta &= \Phi_c(\mathbf{h}_{t-1}, \mathbf{V}), \\ \alpha &= \Phi_s(\mathbf{h}_{t-1}, f_c(\mathbf{V}, \beta)), \\ \mathbf{X} &= f(\mathbf{V}, \alpha, \beta).\end{aligned}\quad (8)$$

where  $f_c(\cdot)$  is a channel-wise multiplication for feature map channels and corresponding channel weights.

**Spatial-Channel.** The second type denoted as Spatial-Channel (S-C) is a model with spatial attention implemented first. For S-C type, given an initial feature map  $\mathbf{V}$ , we first utilize spatial attention  $\Phi_s$  to obtain the spatial attention weights  $\alpha$ . Based on  $\alpha$ , the linear function  $f_s(\cdot)$ , and the channel-wise attention model  $\Phi_c$ , we can calculate the modulated feature  $\mathbf{X}$  following the recipe of C-S type:

$$\begin{aligned}\alpha &= \Phi_s(\mathbf{h}_{t-1}, \mathbf{V}), \\ \beta &= \Phi_c(\mathbf{h}_{t-1}, f_s(\mathbf{V}, \alpha)), \\ \mathbf{X} &= f(\mathbf{V}, \alpha, \beta).\end{aligned}\quad (9)$$

where  $f_s(\cdot)$  is an element-wise multiplication for regions of each feature map channel and its corresponding region attention weights.

## 4. Experiments

We will validate the effectiveness of the proposed SCA-CNN framework for image captioning by answering the following questions: **Q1** Is the channel-wise attention effective? Will it improve the spatial attention? **Q2** Is the multi-

layer attention effective? **Q3** How does SCA-CNN perform compared to other state-of-the-art visual attention models?

#### 4.1. Dataset and Metric

We conducted experiments on three well-known benchmarks: 1) **Flickr8k** [10]: it contains 8,000 images. According to its official split, it selects 6,000 images for training, 1,000 images for validation, and 1,000 images for testing; 2) **Flickr30k** [38]: it contains 31,000 images. Because of the lack of official split, for fair comparison with previous works, we reported results in a publicly available split used in previous work [13]. In this split, 29,000 images are used for training, 1,000 images for validation, and 1,000 images for testing; and 3) **MSCOCO** [16]: it contains 82,783 images in training set, 40,504 images in validation set and 40,775 images in test set. As the ground truth of MSCOCO test set is not available, the validation set is further splitted into a validation subset for model selection and a test subset for local experiments. This split also follows [13]. It utilizes the whole 82,783 training set images for training, and selects 5,000 images for validation and 5,000 images for test from official validation set. As for the sentences preprocessing, we followed the publicly available code<sup>1</sup>. We used BLEU (**B@1,B@2,B@3,B@4**) [20], METEOR (**MT**) [3], CIDEr(**CD**) [28], and ROUGE-L (**RG**) [15] as evaluation metrics. For all the four metrics, in a nutshell, they measure the consistency between n-gram occurrences in generated sentences and ground-truth sentences, where this consistency is weighted by n-gram saliency and rarity. Meanwhile, all the four metrics can be calculated directly through the MSCOCO caption evaluation tool<sup>2</sup>. And our source code is already publicly available<sup>3</sup>.

#### 4.2. Setup

In our captioning system, for image encoding part, we adopted two widely-used CNN architectures: VGG-19 [25] and ResNet-152 [8] as the basic CNNs for SCA-CNN. For the caption decoding part, we used an LSTM [9] to generate caption words. Word embedding dimension and LSTM hidden state dimension are respectively set to 100 and 1,000. The common space dimension for calculating attention weights is set to 512 for both two type attention. For Flickr8k, mini-batch size is set to 16, and for Flickr30k and MSCOCO, mini-batch size is set to 64. We use dropout and early stopping to avoid overfitting. Our whole framework is trained in an end-to-end way with Adadelta [39], which is a stochastic gradient descent method using an adaptive learning rate algorithm. The caption generation process would be halted until a special END token is predicted or a predefined max sentence length is reached. We followed the strategy

of BeamSearch [31] in the testing period, which selects the best caption from some candidates, and the beam size is set to 5. We noticed a trick that incorporates beam search with length normalization [11] which can help to improve performance in some degree. But for fair comparisons, all results reported are without length normalization.

#### 4.3. Evaluations of Channel-wise Attention (Q1)

**Comparing Methods.** We first compared spatial attention with channel-wise attention. 1) **S**: It is a pure spatial attention model. After obtaining spatial attention weights based on the last conv-layer, we use element-wise multiplication to produce a spatial weighted feature. For VGG-19 and ResNet-152, the last conv-layer represents *conv5\_4* layer and *res5c*, respectively. Instead of regarding the weighted feature map as the final visual representation, we feed the spatial weighted feature into their own following CNN layers. For VGG-19, there are two fully-connected layers follows *conv5\_4* layer and for ResNet-152, *res5c* layer is followed by a mean pooling layer. 2) **C**: It is a pure channel-wise attention model. The whole strategy for the C type model is same as S type. The only difference is substituting the spatial attention with channel-wise attention as Eq. (4). 3) **C-S**: This is the first type model incorporating two attention mechanisms as Eq. (8). 4) **S-C**: Another incorporating model introduced in Eq. (9). 5) **SAT**: It is the “hard” attention model introduced in [34]. The reason why we report the results of “hard” attention instead of the “soft” attention is that “hard” attention always has better performance on different datasets and metrics. SAT is also a pure spatial attention model like S. But there are two main differences. The first one is the strategy of modulating visual feature with attention weights. The second one is whether to feed the attending features into their following layers. All VGG results reported in Table 1 came from the original paper and ResNet results are our own implementation.

**Results** From Table 1, we have the following observations: 1) For VGG-19, performance of S is better than that of SAT; but for ResNet-152, the results are opposite. This is because the VGG-19 network has fully-connected layers, which can preserve spatial information. Instead, in ResNet-152, the last conv-layer is originally followed by an average pooling layer, which can destroy spatial information. 2) Comparing to the performance of S, the performance of C can be significant improved in ResNet-152 rather than VGG-19. It shows that the more channel numbers can help improve channel-wise attention performance in the sense that ResNet-152 has more channel numbers (*i.e.* 2048) than VGG-19 (*i.e.* 512). 3) In ResNet-152, both C-S and S-C can achieve better performance than S. This demonstrates that we can improve performance significantly by adding channel-wise attention as long as channel numbers are large. 4) In both of two networks, the per-

<sup>1</sup><https://github.com/karpathy/neuraltalk>

<sup>2</sup><https://github.com/tylin/coco-caption>

<sup>3</sup><https://github.com/zjuchenlong/sca-cnn>

Dataset	Network	Method	B@4	MT	RG	CD
Flickr8k	VGG	S	23.0	21.0	49.1	<b>60.6</b>
		SAT	21.3	20.3	—	—
		C	22.6	20.3	48.7	58.7
		S-C	22.6	20.9	48.7	<b>60.6</b>
		C-S	<b>23.5</b>	<b>21.1</b>	<b>49.2</b>	60.3
Flickr30k	ResNet	S	20.5	19.6	47.4	49.9
		SAT	21.7	20.1	48.4	55.5
		C	24.4	21.5	50.0	65.5
		S-C	24.8	<b>22.2</b>	50.5	65.1
		C-S	<b>25.7</b>	22.1	<b>50.9</b>	<b>66.5</b>
MS COCO	VGG	S	<b>21.1</b>	18.4	43.1	<b>39.5</b>
		SAT	19.9	<b>18.5</b>	—	—
		C	20.1	18.0	42.7	38.0
		S-C	20.8	17.8	42.9	38.2
		C-S	21.0	18.0	<b>43.3</b>	38.5
MS COCO	ResNet	S	20.5	17.4	42.8	35.3
		SAT	20.1	17.8	42.9	36.3
		C	21.5	18.4	43.8	42.2
		S-C	21.9	18.5	44.0	<b>43.1</b>
		C-S	<b>22.1</b>	<b>19.0</b>	<b>44.6</b>	42.5

Table 1. The performance of S, C, C-S, S-C, SAT with one attentive layer in VGG-19 and ResNet-152.

formance of S-C and C-S is quite close. Generally, C-S is slightly better than S-C, so in the following experiments we use C-S to represent incorporating model.

#### 4.4. Evaluations of Multi-layer Attention (Q2)

**Comparing Methods** We will investigate whether we can improve the spatial attention or channel-wise attention performance by adding more attentive layers. We conduct ablation experiments about different number of attentive layer in S and C-S models. In particular, we denote **1-layer**, **2-layer**, **3-layer** as the number of layers equipped with attention, respectively. For VGG-19, 1-st layer, 2-nd layer, 3-rd layer represent *conv5\_4*, *conv5\_3*, *conv5\_2* conv-layer, respectively. As for ResNet-152, it represents *res5c*, *res5c\_branch2b*, *res5c\_branch2a* conv-layer. Specifically, our strategy for training more attentive layers model is to utilize previous trained attentive layer weights as initialization, which can significantly reduce the training time and achieve better results than randomly initialized.

**Results** From Table 2 and 3, we have following obser-

Dataset	Network	Method	B@4	MT	RG	CD
Flickr8k	VGG	1-layer	<b>23.0</b>	21.0	<b>49.1</b>	<b>60.6</b>
		2-layer	22.8	<b>21.2</b>	49.0	60.4
		3-layer	21.6	20.9	48.4	54.5
Flickr30k	ResNet	1-layer	20.5	19.6	47.4	49.9
		2-layer	22.9	21.2	48.8	58.8
		3-layer	<b>23.9</b>	<b>21.3</b>	<b>49.7</b>	<b>61.7</b>
MS COCO	VGG	1-layer	21.1	18.4	43.1	<b>39.5</b>
		2-layer	<b>21.9</b>	<b>18.5</b>	<b>44.3</b>	<b>39.5</b>
		3-layer	20.8	18.0	43.0	38.5
MS COCO	ResNet	1-layer	20.5	17.4	42.8	35.3
		2-layer	20.6	18.6	43.2	39.7
		3-layer	<b>21.0</b>	<b>19.2</b>	<b>43.4</b>	<b>43.5</b>

Table 2. The performance of multi-layer in S in both VGG-19 network and ResNet-152 network

Dataset	Network	Method	B@4	MT	RG	CD
Flickr8k	VGG	1-layer	<b>23.5</b>	21.1	49.2	60.3
		2-layers	22.8	<b>21.6</b>	<b>49.5</b>	62.1
		3-layers	22.7	21.3	49.3	<b>62.3</b>
Flickr30k	ResNet	1-layer	25.7	22.1	50.9	66.5
		2-layers	<b>25.8</b>	22.4	<b>51.3</b>	67.1
		3-layers	25.3	<b>22.9</b>	51.2	<b>67.5</b>
MS COCO	VGG	1-layer	21.0	18.0	43.3	38.5
		2-layers	<b>21.8</b>	<b>18.8</b>	<b>43.7</b>	<b>41.4</b>
		3-layers	20.7	18.3	43.6	39.2
MS COCO	ResNet	1-layer	22.1	19.0	44.6	42.5
		2-layers	<b>22.3</b>	<b>19.5</b>	<b>44.9</b>	<b>44.7</b>
		3-layers	22.0	19.2	44.7	42.8

Table 3. The performance of multi-layer in C-S in both VGG-19 network and ResNet-152 network

vations: 1) In most experiments, adding more attentive layers can achieve better results among two models. The reason is that applying an attention mechanism in multi-layer can help gain visual attention on multiple level semantic abstractions. 2) Too many layers are also prone to resulting in severe overfitting. For example, Flickr8k’s performance is easier to degrade than MSCOCO when adding more attentive layers, as the size of train set of Flickr8k (*i.e.* 6,000) is much smaller than that of MSCOCO (*i.e.* 82,783).

Model	Flickr8k					Flickr30k					MS COCO				
	B@1	B@2	B@3	B@4	MT	B@1	B@2	B@3	B@4	MT	B@1	B@2	B@3	B@4	MT
Deep VS [13]	57.9	38.3	24.5	16.0	–	57.3	36.9	24.0	15.7	–	62.5	45.0	32.1	23.0	19.5
Google NIC [31] <sup>†</sup>	63.0	41.0	27.0	–	–	66.3	42.3	27.7	18.3	–	66.6	46.1	32.9	24.6	–
m-RNN [18]	–	–	–	–	–	60.0	41.0	28.0	19.0	–	67.0	49.0	35.0	25.0	–
Soft-Attention [34]	67.0	44.8	29.9	19.5	18.9	66.7	43.4	28.8	19.1	18.5	70.7	49.2	34.4	24.3	23.9
Hard-Attention [34]	67.0	45.7	31.4	21.3	20.3	<b>66.9</b>	43.9	29.6	19.9	18.5	71.8	50.4	35.7	25.0	23.0
emb-gLSTM [11]	64.7	45.9	31.8	21.2	20.6	64.6	44.6	30.5	20.6	17.9	67.0	49.1	35.8	26.4	22.7
ATT [37] <sup>†</sup>	–	–	–	–	–	64.7	46.0	32.4	<b>23.0</b>	18.9	70.9	53.7	40.2	30.4	24.3
SCA-CNN-VGG	65.5	46.6	32.6	22.8	21.6	64.6	45.3	31.7	21.8	18.8	70.5	53.3	39.7	29.8	24.2
SCA-CNN-ResNet	<b>68.2</b>	<b>49.6</b>	<b>35.9</b>	<b>25.8</b>	<b>22.4</b>	66.2	<b>46.8</b>	<b>32.5</b>	22.3	<b>19.5</b>	<b>71.9</b>	<b>54.8</b>	<b>41.1</b>	<b>31.1</b>	<b>25.0</b>

Table 4. Performances compared with the state-of-art in Flickr8k, Flickr30k and MSCOCO dataset. SCA-CNN-VGG is our C-S 2-layer model based on VGG-19 network, and SCA-CNN-ResNet is our C-S 2-layer model based on ResNet-152 network. <sup>†</sup> indicates an ensemble model results. (–) indicates an unknown metric

Model	B@1		B@2		B@3		B@4		METEOR		ROUGE-L		CIDEr	
	c5	c40	c5	c40	c5	c40								
SCA-CNN	71.2	89.4	54.2	80.2	40.4	69.1	30.2	57.9	24.4	33.1	52.4	67.4	91.2	92.1
Hard-Attention	70.5	88.1	52.8	77.9	38.3	65.8	27.7	53.7	24.1	32.2	51.6	65.4	86.5	89.3
ATT <sup>†</sup>	<b>73.1</b>	<b>90.0</b>	<b>56.5</b>	<b>81.5</b>	<b>42.4</b>	<b>70.9</b>	<b>31.6</b>	<b>59.9</b>	25.0	<b>33.5</b>	53.5	<b>68.2</b>	<b>95.3</b>	<b>95.8</b>
Google NIC <sup>†</sup>	71.3	89.5	54.2	80.2	40.7	69.4	30.9	58.7	<b>25.4</b>	<b>34.6</b>	53.0	<b>68.2</b>	94.3	94.6

Table 5. Performances of the proposed attention model on the onlines MSCOCO testing server. <sup>†</sup> indicates an ensemble model results.

#### 4.5. Comparison with State-of-The-Arts (Q3)

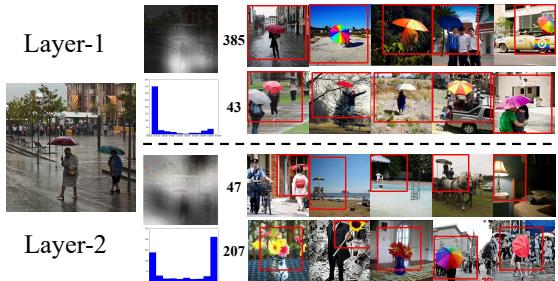
**Comparing Methods** We compared the proposed SCA-CNN with state-of-the-art image captioning models. 1) **Deep VS** [13], **m-RNN** [18], and **Google NIC** [31] are all end-to-end multimodal networks, which combine CNNs for image encoding and RNN for sequence modeling. 2) **Soft-Attention** [34] and **Hard-Attention** [34] are both pure spatial attention model. The “soft” attention weighted sums up the visual features as the attending feature, while the “hard” one randomly samples the region feature as the attending feature. 3) **emb-gLSTM** [11] and **ATT** [37] are both semantic attention models. For emb-gLSTM, it utilizes correlation between image and its description as gloabl semantic information, and for ATT it utilizes visual concepts corresponded words as semantic information. The results reported in Table 4 are from the 2-layer C-S model for both VGG-19 and ResNet-152 network, since this type model always obtains the best performance in previous experiments. Besides the three benchmarks, we also evaluated our model on MSCOCO Image Challenge set c5 and c40 by uploading results to the official test sever. The results are reported in Table 5.

**Results** From Table 4 and Table 5, we can see that in most cases, SCA-CNN outperforms the other models. This is due to the fact that SCA-CNN exploits spatial, channel-wise, and multi-layer attentions, while most of other attention models only consider one attention type. The rea-

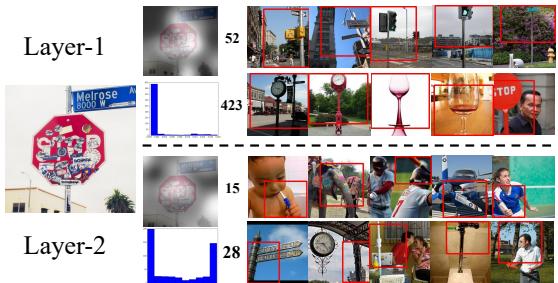
sons why we cannot surpass ATT and Google NIC come from two sides: 1) Both ATT and Google NIC use ensemble models, while SCA-CNN is a single model; ensemble models can always obtain better results than single one. 2) More advanced CNN architectures are used; as Google NIC adopts Inception-v3 [27] which has a better classification performance than ResNet which we adopted. In local experiments, on the MSCOCO dataset, ATT surpasses SCA-CNN only 0.6% in BLEU4 and 0.1% in METEOR, respectively. For the MSCOCO server results, Google NIC surpass SCA-CNN only 0.7% in BLEU4 and 1% in METEOR, respectively.

#### 4.6. Visualization of Spatial and Channel-wise Attention

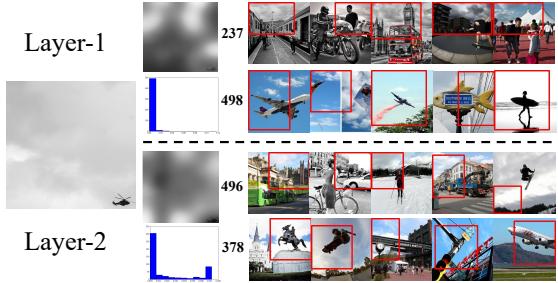
We provided some qualitative examples in Figure 3 for a better understanding of our model. For simplicity, we only visualized results at one word prediction step. For example in the first sample, when SCA-CNN model tries to predict word `umbrella`, our channel-wise attention will assign more weights on feature map channels generated by filters according to the semantics like umbrella, stick, and round-like shape. The histogram in each layer indicates the probability distribution of all channels. The map above histogram is the spatial attention map and white indicates the spatial regions where the model roughly attends to. For each layer we selected two channels with highest channel-wise attention probability. To show the semantic information of



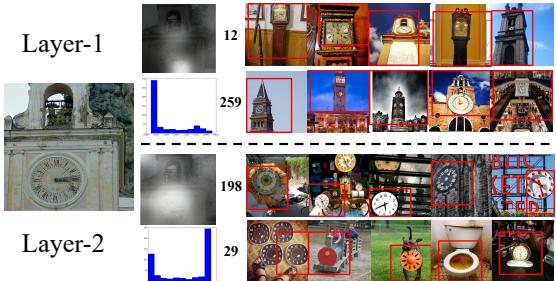
Ours: a woman walking down a street holding an umbrella  
 SAT: a group of people standing next to each other  
 GT: two females walking in the rain with umbrellas



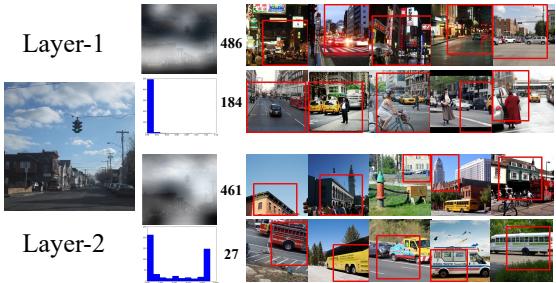
Ours: a street sign on a pole in front of a building  
 SAT: a street sign in front of a building  
 GT: a stop sign is covered with stickers and graffiti



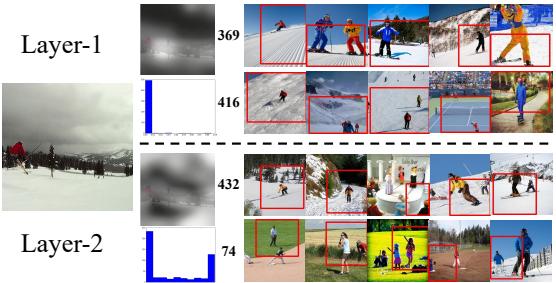
Ours: a plane flying in the sky over a cloudy sky  
 SAT: a plane flying through the sky in the sky  
 GT: a couple of helicopters are in the sky



Ours: a clock tower in the middle of a city  
 SAT: a clock tower on the side of a building  
 GT: there is an old clock on top of a bell tower



Ours: a traffic light in the middle of a city street  
 SAT: a group of people walking down a street  
 GT: a street light at an intersection in a small town



Ours: a man riding skis down a snow covered slope  
 SAT: a man riding a snowboard down a snowy hill  
 GT: a person riding skis goes down a snowy path

Figure 3. Examples of visualization results on spatial attention and channel-wise attention. Each example contains three captions. Ours(SCA-CNN), SAT(hard-attention) and GT(ground truth). The numbers in the third column are the channel numbers of VGG-19 network with highest channel attention weights, and next five images are selected from MSCOCO train set with high activation in the corresponding channel. The red boxes are respective fields in their corresponding layers

the corresponding CNN filter, we used the same methods in [40]. And the red boxes indicate their respective fields.

## 5. Conclusions

In this paper, we proposed a novel deep attention model dubbed SCA-CNN for image captioning. SCA-CNN takes full advantage of characteristics of CNN to yield attentive

image features: spatial, channel-wise, and multi-layer, thus achieving state-of-the-art performance on popular benchmarks. The contribution of SCA-CNN is not only the more powerful attention model, but also a better understanding of where (*i.e.*, spatial) and what (*i.e.*, channel-wise) the attention looks like in a CNN that evolves during sentence generation. In future work, we intend to bring temporal attention

in SCA-CNN, in order to attend features in different video frames for video captioning. We will also investigate how to increase the number of attentive layers without overfitting.

## References

- [1] S. Antol, A. Agrawal, J. Lu, M. Mitchell, D. Batra, C. Lawrence Zitnick, and D. Parikh. Vqa: Visual question answering. In *ICCV*, 2015. 2
- [2] D. Bahdanau, K. Cho, and Y. Bengio. Neural machine translation by jointly learning to align and translate. In *ICLR*, 2014. 2
- [3] S. Banerjee and A. Lavie. Meteor: An automatic metric for mt evaluation with improved correlation with human judgments. In *ACL*, 2005. 5
- [4] K. Chen, J. Wang, L.-C. Chen, H. Gao, W. Xu, and R. Nevatia. Abc-cnn: An attention based convolutional neural network for visual question answering. In *CVPR*, 2016. 1
- [5] M. Corbetta and G. L. Shulman. Control of goal-directed and stimulus-driven attention in the brain. *Nature reviews neuroscience*, 2002. 1
- [6] J. Donahue, L. Anne Hendricks, S. Guadarrama, M. Rohrbach, S. Venugopalan, K. Saenko, and T. Darrell. Long-term recurrent convolutional networks for visual recognition and description. In *CVPR*, 2015. 2
- [7] H. Gao, J. Mao, J. Zhou, Z. Huang, L. Wang, and W. Xu. Are you talking to a machine? dataset and methods for multilingual image question. In *NIPS*, 2015. 2
- [8] K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition. 2016. 1, 2, 3, 5
- [9] S. Hochreiter and J. Schmidhuber. Long short-term memory. *Neural computation*, 1997. 2, 5
- [10] M. Hodosh, P. Young, and J. Hockenmaier. Framing image description as a ranking task: Data, models and evaluation metrics. *JAIR*, 2013. 5
- [11] X. Jia, E. Gavves, B. Fernando, and T. Tuytelaars. Guiding the long-short term memory model for image caption generation. In *ICCV*, 2015. 2, 5, 7
- [12] X. Jiang, F. Wu, X. Li, Z. Zhao, W. Lu, S. Tang, and Y. Zhuang. Deep compositional cross-modal learning to rank via local-global alignment. In *ACM MM*, pages 69–78, 2015. 2
- [13] A. Karpathy and L. Fei-Fei. Deep visual-semantic alignments for generating image descriptions. In *CVPR*, 2015. 2, 5, 7
- [14] R. Krishna, Y. Zhu, O. Groth, J. Johnson, K. Hata, J. Kravitz, S. Chen, Y. Kalantidis, L.-J. Li, D. A. Shamma, et al. Visual genome: Connecting language and vision using crowdsourced dense image annotations. *IJCV*, 2016. 2
- [15] C.-Y. Lin. Rouge: A package for automatic evaluation of summaries. In *ACL*, 2004. 5
- [16] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick. Microsoft coco: Common objects in context. In *ECCV*, 2014. 5
- [17] M. Malinowski, M. Rohrbach, and M. Fritz. Ask your neurons: A neural-based approach to answering questions about images. In *ICCV*, 2015. 2
- [18] J. Mao, W. Xu, Y. Yang, J. Wang, Z. Huang, and A. Yuille. Deep captioning with multimodal recurrent neural networks (m-rnn). In *ICLR*, 2015. 7
- [19] V. Mnih, N. Heess, A. Graves, et al. Recurrent models of visual attention. In *NIPS*, 2014. 1
- [20] K. Papineni, S. Roukos, T. Ward, and W.-J. Zhu. Bleu: a method for automatic evaluation of machine translation. In *ACL*, 2002. 5
- [21] M. Ren, R. Kiros, and R. Zemel. Exploring models and data for image question answering. In *NIPS*, 2015. 2
- [22] P. H. Seo, Z. Lin, S. Cohen, X. Shen, and B. Han. Hierarchical attention networks. *arXiv preprint arXiv:1606.02393*, 2016. 2
- [23] F. Shen, C. Shen, W. Liu, and H. Tao Shen. Supervised discrete hashing. In *CVPR*, pages 37–45, 2015. 2
- [24] F. Shen, C. Shen, Q. Shi, A. Van Den Hengel, and Z. Tang. Inductive hashing on manifolds. In *CVPR*, pages 1562–1569, 2013. 2
- [25] K. Simonyan and A. Zisserman. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*, 2014. 1, 2, 3, 5
- [26] M. F. Stollenga, J. Masci, F. Gomez, and J. Schmidhuber. Deep networks with internal selective attention through feedback connections. In *NIPS*, 2014. 1
- [27] C. Szegedy, V. Vanhoucke, S. Ioffe, J. Shlens, and Z. Wojna. Rethinking the inception architecture for computer vision. In *CVPR*, pages 2818–2826, 2016. 7
- [28] R. Vedantam, C. Lawrence Zitnick, and D. Parikh. Cider: Consensus-based image description evaluation. In *CVPR*, 2015. 5
- [29] S. Venugopalan, M. Rohrbach, J. Donahue, R. Mooney, T. Darrell, and K. Saenko. Sequence to sequence-video to text. In *ICCV*, 2015. 2
- [30] S. Venugopalan, H. Xu, J. Donahue, M. Rohrbach, R. Mooney, and K. Saenko. Translating videos to natural language using deep recurrent neural networks. In *NAACL-HLT*, 2015. 2
- [31] O. Vinyals, A. Toshev, S. Bengio, and D. Erhan. Show and tell: A neural image caption generator. In *CVPR*, 2015. 2, 5, 7
- [32] Y. Wei, W. Xia, M. Lin, J. Huang, B. Ni, J. Dong, Y. Zhao, and S. Yan. Hcp: A flexible cnn framework for multi-label image classification. *TPAMI*, 2016. 1
- [33] H. Xu and K. Saenko. Ask, attend and answer: Exploring question-guided spatial attention for visual question answering. In *ECCV*, 2016. 1, 2
- [34] K. Xu, J. Ba, R. Kiros, K. Cho, A. Courville, R. Salakhutdinov, R. S. Zemel, and Y. Bengio. Show, attend and tell: Neural image caption generation with visual attention. In *ICML*, 2015. 1, 2, 3, 5, 7
- [35] Z. Yang, X. He, J. Gao, L. Deng, and A. Smola. Stacked attention networks for image question answering. In *CVPR*, 2016. 1, 2
- [36] L. Yao, A. Torabi, K. Cho, N. Ballas, C. Pal, H. Larochelle, and A. Courville. Describing videos by exploiting temporal structure. In *ICCV*, 2015. 1
- [37] Q. You, H. Jin, Z. Wang, C. Fang, and J. Luo. Image captioning with semantic attention. In *CVPR*, 2016. 2, 7
- [38] P. Young, A. Lai, M. Hodosh, and J. Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *TACL*, 2014. 5
- [39] M. D. Zeiler. Adadelta: an adaptive learning rate method. *arXiv preprint arXiv:1212.5701*, 2012. 5
- [40] M. D. Zeiler and R. Fergus. Visualizing and understanding convolutional networks. In *ECCV*, 2014. 1, 2, 8
- [41] H. Zhang, Z. Kyaw, S.-F. Chang, and T.-S. Chua. Visual translation embedding network for visual relation detection. In *CVPR*, 2017. 2
- [42] Z. Zhao, H. Lu, C. Deng, X. He, and Y. Zhuang. Partial multi-modal sparse coding via adaptive similarity structure regularization. In *ACM MM*, pages 152–156, 2016. 2
- [43] Y. Zhu, O. Groth, M. Bernstein, and L. Fei-Fei. Visual7w: Grounded question answering in images. In *CVPR*, 2016. 2